

# HousePricing

February 8, 2021

## 1 Housing Prices: Can you predict the home sales prices in Melbourne?

### 1.0.1 Initialize

```
[1]: import math
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### 1.0.2 Load Data

```
[2]: dfMlb = pd.read_csv('house_prices.csv')
dfMlb.head()
```

```
[2]:
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	\
0	Abbotsford	85 Turner St	2	h	1480000	S	Biggin	03/12/16	
1	Abbotsford	25 Bloomburg St	2	h	1035000	S	Biggin	04/02/16	
2	Abbotsford	5 Charles St	3	h	1465000	SP	Biggin	04/03/17	
3	Abbotsford	40 Federation La	3	h	850000	PI	Biggin	04/03/17	
4	Abbotsford	55a Park St	4	h	1600000	VB	Nelson	04/06/16	

	Distance	Postcode	...	Bathroom	Car	Landsize	BuildingArea	YearBuilt	\
0	2.5	3067	...	1	1.0	202	NaN	NaN	
1	2.5	3067	...	1	0.0	156	79.0	1900.0	
2	2.5	3067	...	2	0.0	134	150.0	1900.0	
3	2.5	3067	...	2	1.0	94	NaN	NaN	
4	2.5	3067	...	1	2.0	120	142.0	2014.0	

	CouncilArea	Lattitude	Longitude	Regionname	Propertycount
0	Yarra	-37.7996	144.9984	Northern Metropolitan	4019
1	Yarra	-37.8079	144.9934	Northern Metropolitan	4019
2	Yarra	-37.8093	144.9944	Northern Metropolitan	4019
3	Yarra	-37.7969	144.9969	Northern Metropolitan	4019
4	Yarra	-37.8072	144.9941	Northern Metropolitan	4019

[5 rows x 21 columns]

```
[3]: dfMlb.columns # Display all columns
```

```
[3]: Index(['Suburb', 'Address', 'Rooms', 'Type', 'Price', 'Method', 'SellerG',  
        'Date', 'Distance', 'Postcode', 'Bedroom2', 'Bathroom', 'Car',  
        'Landsize', 'BuildingArea', 'YearBuilt', 'CouncilArea', 'Lattitude',  
        'Longtitude', 'Regionname', 'Propertycount'],  
        dtype='object')
```

### 1.0.3 Clean Data

```
[4]: # Fill empty cells with the median from that column  
median_year_built = math.floor(dfMlb.YearBuilt.median())  
dfMlb.YearBuilt = dfMlb.YearBuilt.fillna(median_year_built)
```

```
[5]: median_car = math.floor(dfMlb.Car.median())  
dfMlb.Car = dfMlb.Car.fillna(median_car)
```

```
[6]: median_building_area = math.floor(dfMlb.BuildingArea.median())  
dfMlb.BuildingArea = dfMlb.BuildingArea.fillna(median_building_area)
```

```
[7]: # Drop remaining rows with empty cells  
for column in dfMlb.columns:  
    dfMlb = dfMlb[dfMlb[str(column)].notna()]  
# Remove data where values are 0  
dfMlb = dfMlb[dfMlb.Landsize != 0]  
dfMlb = dfMlb[dfMlb.Price != 0]  
# Get average price to calculate MAE in a future  
avg_price = math.floor(dfMlb.Price.mean())
```

```
[8]: # Drop unique, redudant and useless columns  
dfMlb = dfMlb.drop(columns=['Address', 'Date', 'Lattitude', 'Longtitude'])  
# Rename misspelled columns inplace  
dfMlb.rename(columns={'SellerG': 'Seller', 'Bedroom2': 'Bedroom', 'Regionname': 'RegionName', 'Propertycount': 'PropertyCount'}, inplace=True)
```

```
[9]: dfMlb.describe(include='all')
```

```
[9]:
```

	Suburb	Rooms	Type	Price	Method	Seller	\
count	10272	10272.000000	10272	1.027200e+04	10272	10272	
unique	305	NaN	3	NaN	5	237	
top	Reservoir	NaN	h	NaN	S	Nelson	
freq	301	NaN	8033	NaN	7017	1239	
mean	NaN	3.068536	NaN	1.142861e+06	NaN	NaN	
std	NaN	0.899147	NaN	6.493972e+05	NaN	NaN	
min	NaN	1.000000	NaN	1.310000e+05	NaN	NaN	

25%	NaN	3.000000	NaN	7.100000e+05	NaN	NaN
50%	NaN	3.000000	NaN	9.675000e+05	NaN	NaN
75%	NaN	4.000000	NaN	1.400000e+06	NaN	NaN
max	NaN	10.000000	NaN	9.000000e+06	NaN	NaN

	Distance	Postcode	Bedroom	Bathroom	Car \
count	10272.000000	10272.000000	10272.000000	10272.000000	10272.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	10.402434	3102.092971	3.038162	1.570678	1.659073
std	5.578720	91.127321	0.916483	0.707189	0.985341
min	0.000000	3000.000000	0.000000	0.000000	0.000000
25%	6.700000	3044.000000	2.000000	1.000000	1.000000
50%	9.850000	3081.000000	3.000000	1.000000	2.000000
75%	13.000000	3146.000000	4.000000	2.000000	2.000000
max	47.400000	3977.000000	20.000000	8.000000	10.000000

	Landsize	BuildingArea	YearBuilt	CouncilArea \
count	10272.000000	10272.000000	10272.000000	10272
unique	NaN	NaN	NaN	33
top	NaN	NaN	NaN	Moreland
freq	NaN	NaN	NaN	1030
mean	647.930199	141.715375	1965.026967	NaN
std	4548.807775	101.695480	30.162665	NaN
min	1.000000	0.000000	1196.000000	NaN
25%	258.000000	126.000000	1960.000000	NaN
50%	516.000000	126.000000	1970.000000	NaN
75%	670.000000	137.000000	1970.000000	NaN
max	433014.000000	6791.000000	2018.000000	NaN

	RegionName	PropertyCount
count	10272	10272.000000
unique	8	NaN
top	Southern Metropolitan	NaN
freq	3303	NaN
mean	NaN	7304.301889
std	NaN	4392.989470
min	NaN	249.000000
25%	NaN	4019.000000
50%	NaN	6482.000000
75%	NaN	9704.000000
max	NaN	21650.000000

```
[10]: # Check if we have succesfully removed all null vars
dfMlb.isnull().sum()
```

```
[10]: Suburb          0
      Rooms          0
      Type           0
      Price          0
      Method         0
      Seller         0
      Distance       0
      Postcode       0
      Bedroom        0
      Bathroom       0
      Car            0
      Landsize       0
      BuildingArea   0
      YearBuilt      0
      CouncilArea    0
      RegionName     0
      PropertyCount  0
      dtype: int64
```

```
[11]: # Transform numerical values that are categorical to object type
dfMlb['YearBuilt'] = dfMlb['YearBuilt'].astype('object',copy=False)
dfMlb['Postcode'] = dfMlb['Postcode'].astype('object',copy=False)
# Encode categorical variables (to numerical)
from sklearn import model_selection, preprocessing
for col in dfMlb.columns:
    if dfMlb[col].dtype == 'object':
        lbl = preprocessing.LabelEncoder()
        lbl.fit(list(dfMlb[col].values))
        dfMlb[col] = lbl.transform(list(dfMlb[col].values))
```

```
[12]: dfMlb.dtypes
```

```
[12]: Suburb          int64
      Rooms          int64
      Type           int64
      Price          int64
      Method         int64
      Seller         int64
      Distance       float64
      Postcode       int64
      Bedroom        int64
      Bathroom       int64
      Car            float64
      Landsize       int64
      BuildingArea   float64
      YearBuilt      int64
      CouncilArea    int64
```

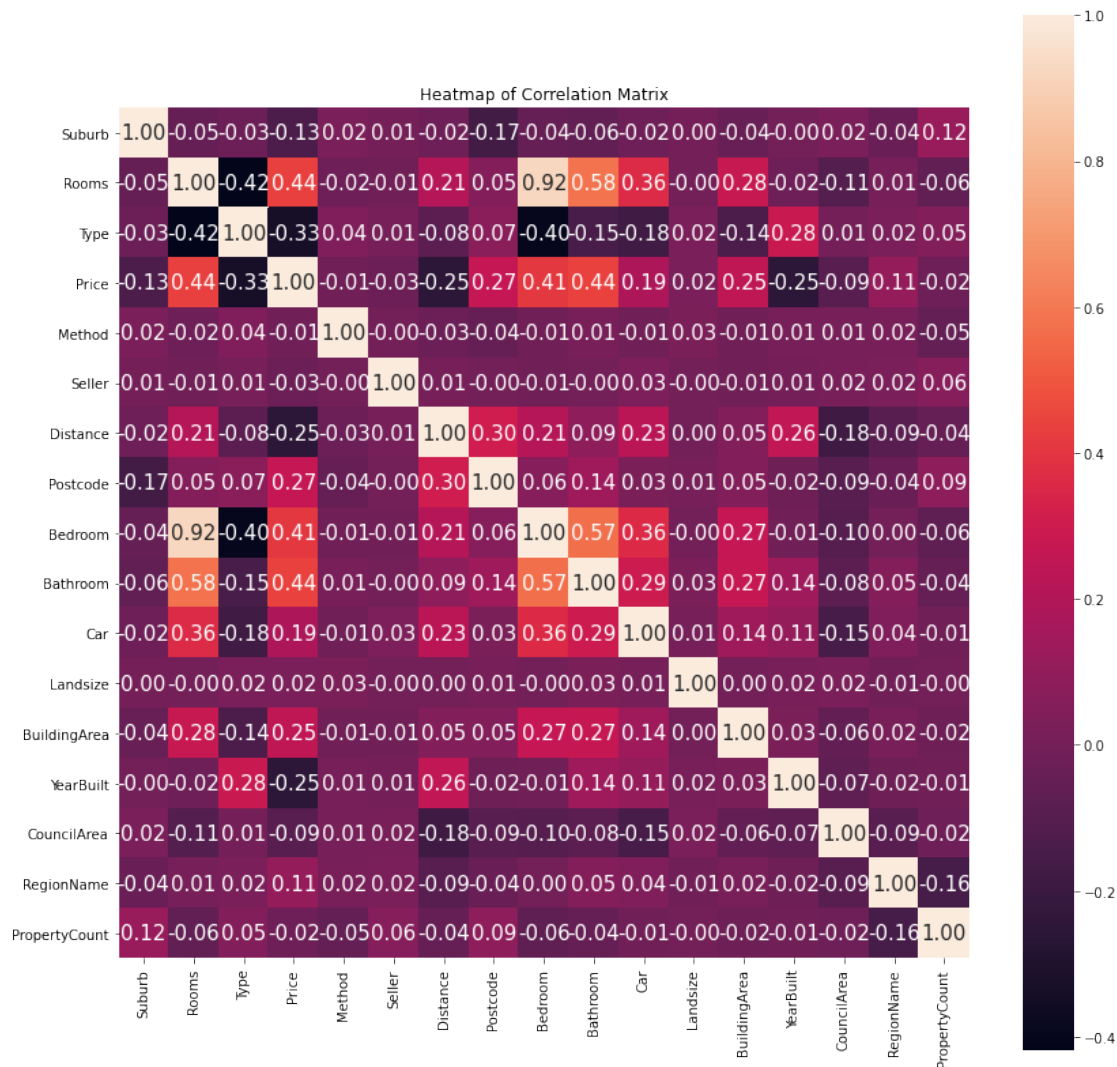
```
RegionName          int64
PropertyCount       int64
dtype: object
```

#### 1.0.4 Feature Engineering

```
[13]: import matplotlib as mpl
      %matplotlib inline
      import warnings
      warnings.filterwarnings('ignore')

      corr = dfMlb.corr()
      corr = (corr)
      plt.figure(figsize=(14,14))
      sns.heatmap(corr, cbar = True, square = True, annot=True, fmt= '.
      ↪2f',annot_kws= {'size': 15}, xticklabels=corr.columns.values,
      ↪yticklabels=corr.columns.values)
      plt.title('Heatmap of Correlation Matrix')
```

```
[13]: Text(0.5, 1.0, 'Heatmap of Correlation Matrix')
```



### 1.0.5 Select Data

```
[14]: # Select a Prediction Target
y = dfMlb.loc[:,['Price']]
```

```
[15]: y.head()
```

```
[15]:      Price
0  1480000
1  1035000
2  1465000
3   850000
4  1600000
```

```
[16]: # Selected Set of Features
x = dfMlb.loc[:, ['Suburb', 'Rooms', 'Type', 'Method', 'Seller', 'Distance',
↳ 'Postcode', 'Bathroom', 'Car', 'Landsize', 'BuildingArea', 'YearBuilt',
↳ 'CouncilArea', 'RegionName', 'PropertyCount']]
```

```
[17]: x.describe()
```

```
[17]:
```

	Suburb	Rooms	Type	Method	Seller \
count	10272.000000	10272.000000	10272.000000	10272.000000	10272.000000
mean	149.215343	3.068536	0.347741	1.362344	115.884930
std	88.614796	0.899147	0.697428	1.077439	69.122205
min	0.000000	1.000000	0.000000	0.000000	0.000000
25%	61.000000	3.000000	0.000000	1.000000	63.000000
50%	150.000000	3.000000	0.000000	1.000000	121.000000
75%	227.000000	4.000000	0.000000	1.000000	173.000000
max	304.000000	10.000000	2.000000	4.000000	236.000000

	Distance	Postcode	Bathroom	Car	Landsize \
count	10272.000000	10272.000000	10272.000000	10272.000000	10272.000000
mean	10.402434	72.466803	1.570678	1.659073	647.930199
std	5.578720	47.399684	0.707189	0.985341	4548.807775
min	0.000000	0.000000	0.000000	0.000000	1.000000
25%	6.700000	33.000000	1.000000	1.000000	258.000000
50%	9.850000	63.000000	1.000000	2.000000	516.000000
75%	13.000000	109.000000	2.000000	2.000000	670.000000
max	47.400000	192.000000	8.000000	10.000000	433014.000000

	BuildingArea	YearBuilt	CouncilArea	RegionName	PropertyCount
count	10272.000000	10272.000000	10272.000000	10272.000000	10272.000000
mean	141.715375	89.670463	14.209502	3.743477	7304.301889
std	101.695480	27.966337	9.796792	2.062906	4392.989470
min	0.000000	0.000000	0.000000	0.000000	249.000000
25%	126.000000	84.000000	6.000000	2.000000	4019.000000
50%	126.000000	94.000000	15.000000	5.000000	6482.000000
75%	137.000000	94.000000	23.000000	5.000000	9704.000000
max	6791.000000	142.000000	32.000000	7.000000	21650.000000

## 1.0.6 Build Model

1. Define
2. Fit
3. Predict
4. Evaluate

```
[18]: # Define
from sklearn.tree import DecisionTreeRegressor
```

```
mdlDtrMlb = DecisionTreeRegressor()
```

```
[19]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳ random_state=1)
```

```
[20]: # Feature Scaling (Normalize IV)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
[21]: # Build decision tree model
from sklearn.metrics import mean_absolute_error

# define
mdlDtrMlbSpl = DecisionTreeRegressor(random_state=1)

# fit
mdlDtrMlbSpl.fit(x_train, y_train)

# predict
y_test_pred = mdlDtrMlbSpl.predict(x_test)

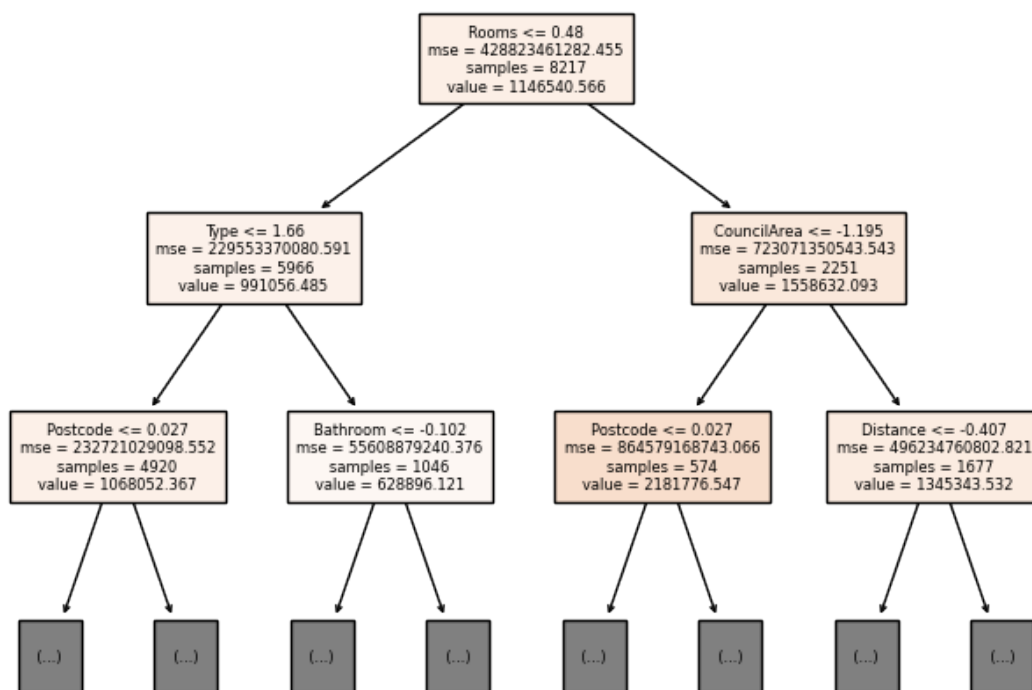
# evaluate
mae = mean_absolute_error(y_test, y_test_pred)
print('MAE (after train-test-split): ', round(mae,2))
print("Current mae is ", round(mae/avg_price*100,2), "%")
```

```
MAE (after train-test-split): 251205.52
Current mae is 21.98 %
```

```
[22]: from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6), dpi=100)
plot_tree(mdlDtrMlbSpl, max_depth=2, feature_names=x.columns, fontsize=6,
↳ filled=True)
plt.show()
```





### 1.0.7 Improve Decision Tree model

```

[23]: # We will try to get our MAE down to 15%
      # Function to tweak max_leaf_nodes parameter
      def get_mae_lfndsDTR(X_trn, X_tst, y_trn, y_tst, mx_lf_nds):
          mdlDtrMlbDpt = DecisionTreeRegressor(random_state=1,
          ↳max_leaf_nodes=mx_lf_nds)
          mdlDtrMlbDpt.fit(X_trn, y_trn)
          y_tst_prd = mdlDtrMlbDpt.predict(X_tst)
          mae = mean_absolute_error(y_test, y_tst_prd)
          return mae

      # Funtion to tweak min_samples_leaf parameter
      def get_mae_smplfsDTR(X_trn, X_tst, y_trn, y_tst, mn_spls_lf):
          mdlDtrMlbDpt = DecisionTreeRegressor(random_state=1,
          ↳min_samples_leaf=mn_spls_lf)
          mdlDtrMlbDpt.fit(X_trn, y_trn)
          y_tst_prd = mdlDtrMlbDpt.predict(X_tst)
          mae = mean_absolute_error(y_test, y_tst_prd)
          return mae
  
```

```
# Funtion to tweak max_depth parameter
def get_mae_mxdpthDTR(X_trn, X_tst, y_trn, y_tst, mx_dpth):
    mdlDtrMlbDpt = DecisionTreeRegressor(random_state=1, max_depth=mx_dpth)
    mdlDtrMlbDpt.fit(X_trn, y_trn)
    y_tst_prd = mdlDtrMlbDpt.predict(X_tst)
    mae = mean_absolute_error(y_test, y_tst_prd)
    return mae
```

```
[24]: # For loops to get the ranges for the hyper parameter tunning
for i in [5, 50, 500, 5000, 50000]:
    mae = get_mae_lfndsDTR(x_train, x_test, y_train, y_test, i)
    print('Max leaf nodes: ', i, '\t MAE', round(mae, 2), '\t % MAE ',
    ↪round(mae/avg_price*100,2))

for i in [5, 10]:
    mae = get_mae_smplfsDTR(x_train, x_test, y_train, y_test, i)
    print('Min samples per leaf: ', i, '\t MAE', round(mae, 2), '\t % MAE ',
    ↪round(mae/avg_price*100,2))

for i in [5, 10, 15, 20]:
    mae = get_mae_mxdpthDTR(x_train, x_test, y_train, y_test, i)
    print('Max depth: ', i, '\t MAE', round(mae, 2), '\t % MAE ', round(mae/
    ↪avg_price*100,2))
```

Max leaf nodes:	5	MAE 338975.62	% MAE 29.66
Max leaf nodes:	50	MAE 247512.97	% MAE 21.66
Max leaf nodes:	500	MAE 226680.59	% MAE 19.83
Max leaf nodes:	5000	MAE 249952.48	% MAE 21.87
Max leaf nodes:	50000	MAE 250348.93	% MAE 21.91
Min samples per leaf:	5	MAE 217297.08	% MAE 19.01
Min samples per leaf:	10	MAE 214211.42	% MAE 18.74
Max depth:	5	MAE 254204.12	% MAE 22.24
Max depth:	10	MAE 219378.22	% MAE 19.2
Max depth:	15	MAE 243619.52	% MAE 21.32
Max depth:	20	MAE 248162.38	% MAE 21.71

## 1.0.8 Hyperparameter tuning for Decision Tree Regressor

```
[25]: mdlDtrMlbSpl.get_params()
```

```
[25]: {'ccp_alpha': 0.0,
      'criterion': 'mse',
      'max_depth': None,
      'max_features': None,
      'max_leaf_nodes': None,
      'min_impurity_decrease': 0.0,
      'min_impurity_split': None,
```

```

'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'presort': 'deprecated',
'random_state': 1,
'splitter': 'best'}

```

```

[ ]: # Import modules for Hyperparameter Tuning
from scipy.stats import randint
from sklearn.model_selection import RandomizedSearchCV

# The function to measure the quality of a split
criterion = ['mse', 'friedman_mse', 'mae']
# The strategy used to choose the split at each node
splitter = ['best', 'random']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(5, 100, num = 2)] + [None]
# Minimum number of samples required at each leaf node
min_samples_leaf = [int(x) for x in range(1, 10)] + [None]

# Setup the parameters and distribution to sample from: param_dist
param_dist = {
    "criterion": criterion,
    "splitter": splitter,
    "max_depth": max_depth,
    "min_samples_leaf": min_samples_leaf
}

# define
mdlDtrMlbDpt = DecisionTreeRegressor(random_state=1)

# Hyper parameter tune
mdlDtrMlbDptCV = RandomizedSearchCV(estimator = mdlDtrMlbDpt,
    ↪ param_distributions = param_dist, cv = 5)

# fit
mdlDtrMlbDptCV.fit(x_train, y_train)

# print the tuned parameters and score
print("Tuned Decision Tree Parameters: {}".format(mdlDtrMlbDptCV.best_params_))
print("Best score is {}", format(mdlDtrMlbDptCV.best_score_))

```

```

[27]: # Predicting the test set results
y_test_pred = mdlDtrMlbDptCV.predict(x_test)

# evaluate
mae = mean_absolute_error(y_test, y_test_pred)

```

```
print('MAE (after train-test-split): ', round(mae,2))
print("Current mae is ", round(mae/avg_price*100,2), "%")
```

MAE (after train-test-split): 231069.11

Current mae is 20.22 %

### 1.0.9 Build Random Forest Model

```
[28]: from sklearn.ensemble import RandomForestRegressor

# define
mdlRfsMlb = RandomForestRegressor(random_state=1)

# fit
mdlRfsMlb.fit(x_train, y_train)

# predict
y_test_pred = mdlRfsMlb.predict(x_test)

# evaluate
mae = mean_absolute_error(y_test, y_test_pred)

print('MAE (Random Forest):', round(mae,2))
print("Current mae is ", round(mae/avg_price*100,2), "%")
```

MAE (Random Forest): 179177.72

Current mae is 15.68 %

### 1.0.10 Improve Random Forest model

```
[29]: # Function to tweak max_leaf_nodes parameter
def get_mae_lfndsRFR(X_trn, X_tst, y_trn, y_tst, mx_lf_nds):
    mdlRfsMlb = RandomForestRegressor(random_state=1, max_leaf_nodes=mx_lf_nds)
    mdlRfsMlb.fit(X_trn, y_trn)
    y_tst_prd = mdlRfsMlb.predict(X_tst)
    mae = mean_absolute_error(y_test, y_tst_prd)
    return mae

# Funtion to tweak min_samples_leaf parameter
def get_mae_smplfsRFR(X_trn, X_tst, y_trn, y_tst, mn_spls_lf):
    mdlRfsMlb = RandomForestRegressor(random_state=1,
    ↪min_samples_leaf=mn_spls_lf)
    mdlRfsMlb.fit(X_trn, y_trn)
    y_tst_prd = mdlRfsMlb.predict(X_tst)
    mae = mean_absolute_error(y_test, y_tst_prd)
    return mae
```

```
# Funtion to tweak min_samples_leaf parameter
def get_mae_dpthRFR(X_trn, X_tst, y_trn, y_tst, mx_dpth):
    mdlRfsMlb = RandomForestRegressor(random_state=1, max_depth=mx_dpth)
    mdlRfsMlb.fit(X_trn, y_trn)
    y_tst_prd = mdlRfsMlb.predict(X_tst)
    mae = mean_absolute_error(y_test, y_tst_prd)
    return mae
```

```
[30]: # Develop similar for loops for the new model
for i in [5, 50, 500, 5000, 50000]:
    mae = get_mae_lfndsRFR(x_train, x_test, y_train, y_test, i)
    print('Max leaf nodes: ', i, '\t MAE', round(mae, 2), '\t % MAE ',
    ↪round(mae/avg_price*100,2))

for i in [5, 10, 15, 20, 25, 30]:
    mae = get_mae_smplfsRFR(x_train, x_test, y_train, y_test, i)
    print('Min samples per leaf: ', i, '\t MAE', round(mae, 2), '\t % MAE ',
    ↪round(mae/avg_price*100,2))

for i in [5, 10, 15, 20]:
    mae = get_mae_dpthRFR(x_train, x_test, y_train, y_test, i)
    print('Max depth: ', i, '\t MAE', round(mae, 2), '\t % MAE ', round(mae/
    ↪avg_price*100,2))
```

Max leaf nodes:	5	MAE 319839.35	% MAE	27.99
Max leaf nodes:	50	MAE 215124.33	% MAE	18.82
Max leaf nodes:	500	MAE 184240.52	% MAE	16.12
Max leaf nodes:	5000	MAE 179938.83	% MAE	15.74
Max leaf nodes:	50000	MAE 179939.17	% MAE	15.74
Min samples per leaf:	5	MAE 176715.67	% MAE	15.46
Min samples per leaf:	10	MAE 180239.45	% MAE	15.77
Min samples per leaf:	15	MAE 183847.22	% MAE	16.09
Min samples per leaf:	20	MAE 188048.31	% MAE	16.45
Min samples per leaf:	25	MAE 192215.27	% MAE	16.82
Min samples per leaf:	30	MAE 194787.86	% MAE	17.04
Max depth:	5	MAE 231046.1	% MAE	20.22
Max depth:	10	MAE 186388.54	% MAE	16.31
Max depth:	15	MAE 178273.73	% MAE	15.6
Max depth:	20	MAE 179051.59	% MAE	15.67

### 1.0.11 Hyperparameter tuning for Random Forest Regressor

```
[31]: mdlRfsMlb.get_params()
```

```
[31]: {'bootstrap': True,
      'ccp_alpha': 0.0,
```

```

'criterion': 'mse',
'max_depth': None,
'max_features': 'auto',
'max_leaf_nodes': None,
'max_samples': None,
'min_impurity_decrease': 0.0,
'min_impurity_split': None,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 100,
'n_jobs': None,
'oob_score': False,
'random_state': 1,
'verbose': 0,
'warm_start': False}

```

```

[32]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# np.linspace returns evenly spaced numbers over a specified interval.
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(5, 50, num = 5)] + [None]
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]

# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}

mdlRfsMlb = RandomForestRegressor()
# Random search of parameters, across 100 different combinations with all
↳ available logical cores
mdlRfsMlbCV = RandomizedSearchCV(estimator = mdlRfsMlb, param_distributions =
↳ random_grid, n_iter = 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
# Fit the random search model
mdlRfsMlbCV.fit(x_train, y_train)

print("Tuned Decision Tree Parameters: {}".format(mdlRfsMlbCV.best_params_))

```

```
print("Best score is {}", format mdlRfsMlbCV.best_score_))
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 2.6min
```

```
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 9.5min
```

```
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 19.6min finished
```

```
Tuned Decision Tree Parameters: {'n_estimators': 1400, 'min_samples_split': 5,
'min_samples_leaf': 1, 'max_features': 'sqrt', 'max_depth': 38, 'bootstrap':
False}
```

```
Best score is {} 0.7907541719586694
```

```
[33]: # Predicting the test set results
y_test_pred = mdlRfsMlbCV.predict(x_test)

# evaluate
mae = mean_absolute_error(y_test, y_test_pred)
print('MAE (after train-test-split): ', round(mae,2))
print("Current mae is ", round(mae/avg_price*100,2), "%")
```

```
MAE (after train-test-split): 174183.14
```

```
Current mae is 15.24 %
```

**1.0.12 We did it! We developed a model with a 15.24% MAE**

**1.0.13 Alejandro Gleason Méndez - ag77698**